



Synchronization

Dr. Jiju Poovancheri
Operating Systems: Recitations
SMU. Winter 2022

<https://computing.iitl.gov/tutorials/pthreads/>

Pthread mutex (pthread.h)

```
pthread_mutex_t mutex;
```

```
/* create and initialize the mutex lock */
```

```
pthread_mutex_init(&mutex, NULL);
```

```
/* acquire the mutex lock */
```

```
pthread_mutex_lock(&mutex);
```

```
//critical section
```

```
/* release the mutex lock */
```

```
pthread_mutex_unlock(&mutex);
```

```
/* Destroy the mutex lock */
```

```
pthread_mutex_destroy(&mutex);
```

POSIX Semaphores (semaphore.h)

sem_init-Initializes a semaphore with a value. Returns 0 if successful, -1 otherwise.

Syntax:

```
int sem_init(sem_t *sem, int pshared, unsigned int value);
```

sem-semaphore object to be initialized

pshared-flag indicating whether sem to be shared

value-value to be initialized to sem

e.g.,

```
sem_t sem;
```

```
sem_init(&sem, 0, 1);
```

POSIX Semaphores (semaphore.h)

sem_wait- Locks a semaphore and returns 0. If the semaphore value is zero 0, the calling process gets blocked. Returns -1 if unsuccessful (deadlock, interrupt etc..)

Syntax:

```
int sem_wait(sem_t *sem);
```

e.g.,

```
sem_t sem;
```

```
sem_wait(&sem);
```

This is the **down(s)** operation

POSIX Semaphores

sem_post- It increments the value of the semaphore and wakes up a blocked process waiting on the semaphore, if any.

Syntax:

```
int sem_post(sem_t *sem);
```

e.g.,

```
sem_t sem;
```

```
sem_post(&sem);
```

This is the **up(s)** operation

POSIX Semaphores

sem_destroy- destroys the semaphore; no threads should be waiting on the semaphore if its destruction is to succeed.

Syntax:

```
int sem_destroy(sem_t *sem);
```

e.g.,

```
sem_t sem;
```

```
sem_destroy(&sem);
```

Semaphore Vs Mutex?

- Signalling Vs Locking mechanisms
- Mutex lock can be unlocked only by the thread that acquired it.
Semaphore value is changed by any thread acquiring or releasing it.
- Single Vs multiple threads

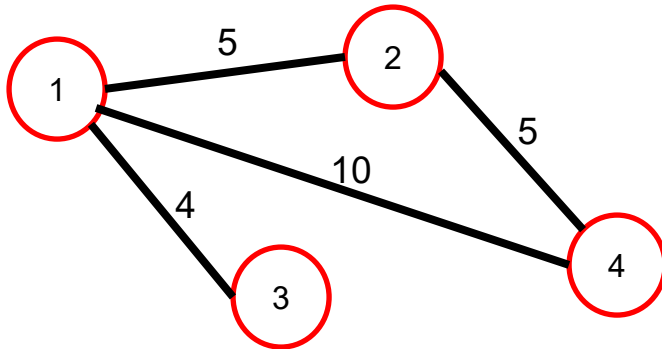
Project 2

Floyd Warshall All Pair Shortest Path

- Given a weighted graph, we want to know the shortest path from one vertex in the graph to another.
 - The Floyd-Warshall algorithm determines the shortest path between all pairs of vertices in a graph.

FWA-Data Structures

- Adjacency matrix (graph)-1 if nodes i and j are connected, else 0.
- Distance matrix (dist)-distance between nodes i and j.
 - dist[i, i]=0;
 - dist[i, j]=INF, if i and j are not connected.



$$\begin{pmatrix}
 0 & 1 & 1 & 1 \\
 1 & 0 & 0 & 1 \\
 1 & 0 & 0 & 0 \\
 1 & 1 & 0 & 0
 \end{pmatrix}$$

graph

$$\begin{pmatrix}
 0 & 5 & 4 & 10 \\
 5 & 0 & \text{INF} & 5 \\
 4 & \text{INF} & 0 & \text{INF} \\
 10 & 5 & \text{INF} & 0
 \end{pmatrix}$$

dist

FW-Single threaded

Algorithm: Floyd-Warshall (FW) algorithm

```
for (int k = 0; k < n; ++k) {  
    for (int i = 0; i < n; ++i) {  
        for (int j = 0; j < n; ++j) {  
            if (dist[i][k] + dist[k][j] < dist[i][j])  
                dist[i][j] = dist[i][k] + dist[k][j]  
        }  
    }  
}
```

FW-Multi-threaded

```
for (int k = 0; k < n; ++k) {  
    for (int i = 0; i < n; ++i) {  
        //set up the arguments to be passed, n, k, i  
        //create threads → pthread_create  
  
    }  
    for (int i = 0; i < n; ++i) {  
        //join threads → pthread_join  
  
    }  
  
}
```

Illustration, D^0



dist[i, j]



dist[k, j]



dist[i, k]

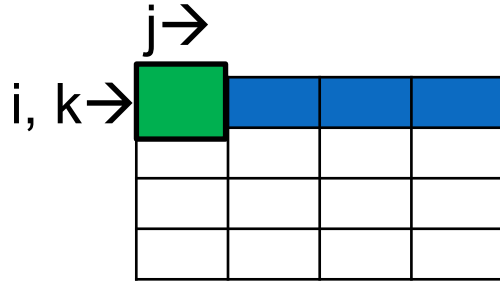
Thread 1

Thread 2

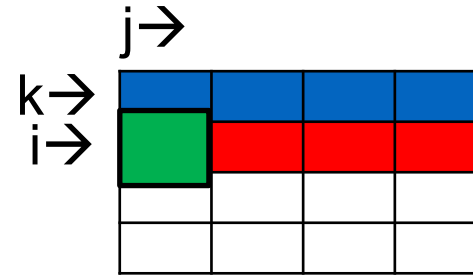
Thread 3

Thread 4

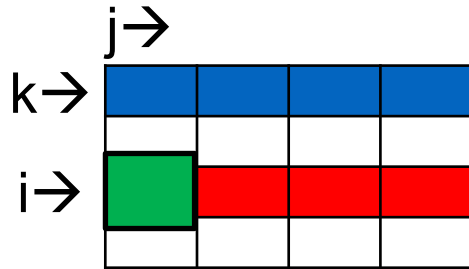
Illustration, D^1 , $k=1$



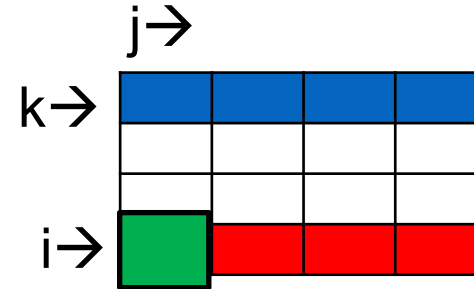
Thread 1



Thread 2

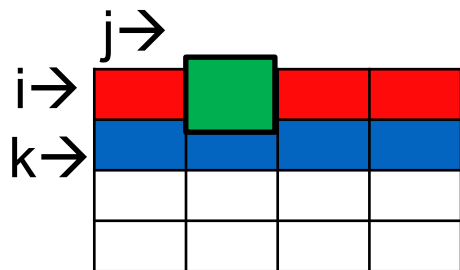


Thread 3

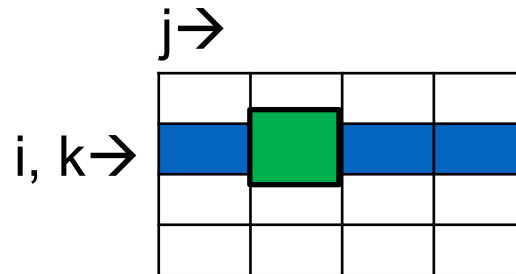


Thread 4

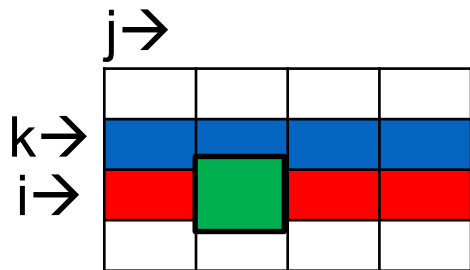
Illustration, $D^2, k=2$



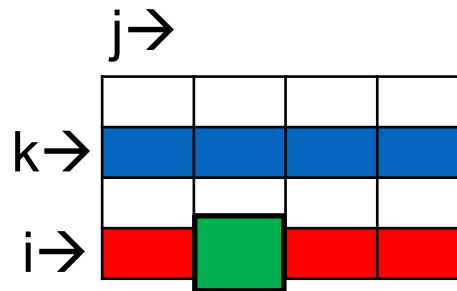
Thread 1



Thread 2

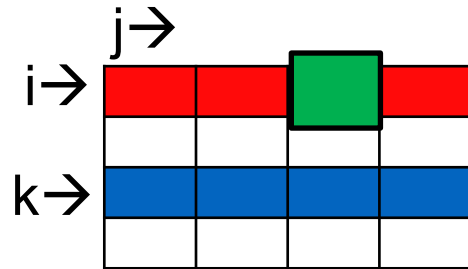


Thread 3

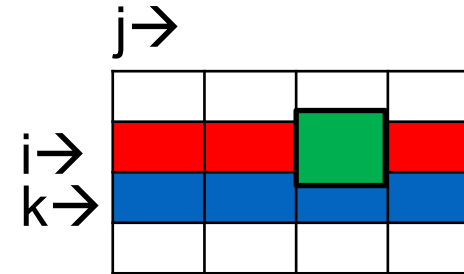


Thread 4

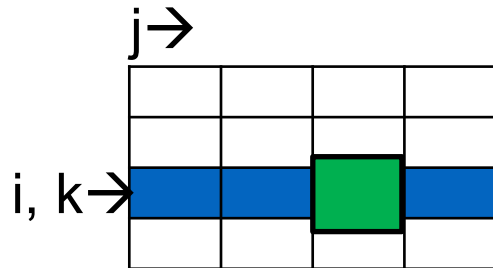
Illustration, D^3 , $k=3$



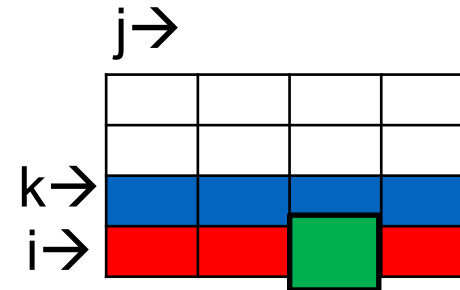
Thread 1



Thread 2

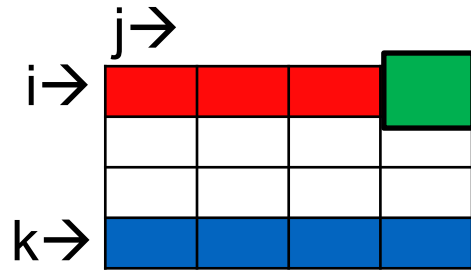


Thread 3

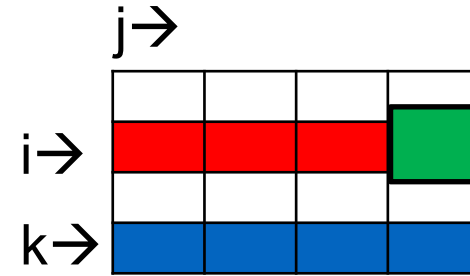


Thread 4

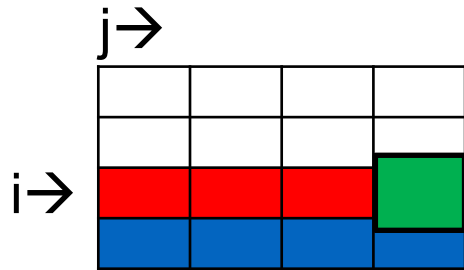
Illustration, D^4 , $k=4$



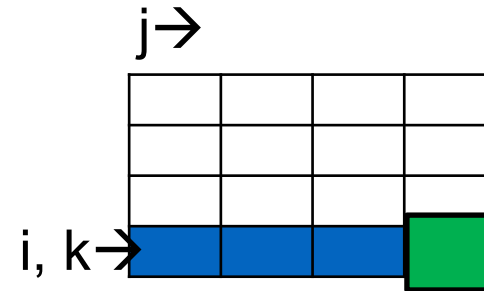
Thread 1



Thread 2



Thread 3



Thread 4

Worker function

```
void *worker(void *args)
{
    //get n,i,k from args
    for (int j = 0; j<n; j++)
    {
        //acquire read lock
        if ((dist[i][k] + dist[k][j]) < dist[i][j])
        {
            //release read lock
            //acquire write lock
            dist[i][j] = dist[i][k] + dist[k][j];
            //release write lock
        }
    }
}
```

Worker function contd..

```
else {  
    //release read lock  
}  
  
}  
pthread_exit(NULL);  
}
```

Passing arguments

- Use structures to hold values n , k and i

e.g.

```
struct arg_s {  
    int n;  
    int i;  
    int k;  
};
```

Passing arguments

- Use structures to hold values n , k and i
- Pass this as arguments when creating pthreads

e.g.

```
int *threads = (pthread_t *)malloc(n *  
sizeof(pthread_t));  
  
for (int i = 0; i < n; ++i)  
pthread_create(&threads[i], NULL, worker, (void  
&arguments[i]));
```

Handling large 2D arrays

```
int **graph; } global  
int **dist;
```

```
//dynamic mem. allocation for 2D array  
graph = malloc(n * sizeof(int *));  
for (int i = 0; i < n; ++i) {  
    graph[i] = malloc(n * sizeof(int))  
}
```